

# Linux Programming

## *5. The File in Context*

Sunwoo Lee

`sunwool@inha.ac.kr`



INHA UNIVERSITY

# Today, we will learn...

- Last week, we learned the basic file handling system calls such as `open()`, `close()`, `read()`, and `write()`.
- Now, we are ready to learn advanced file handling system calls!
  - Ownership
  - File sharing
  - Hard and symbolic link

- ***Files in Multi-user Environment***

- Files with Multiple Names
- Obtaining File Information

# Users and Ownership (1/3)

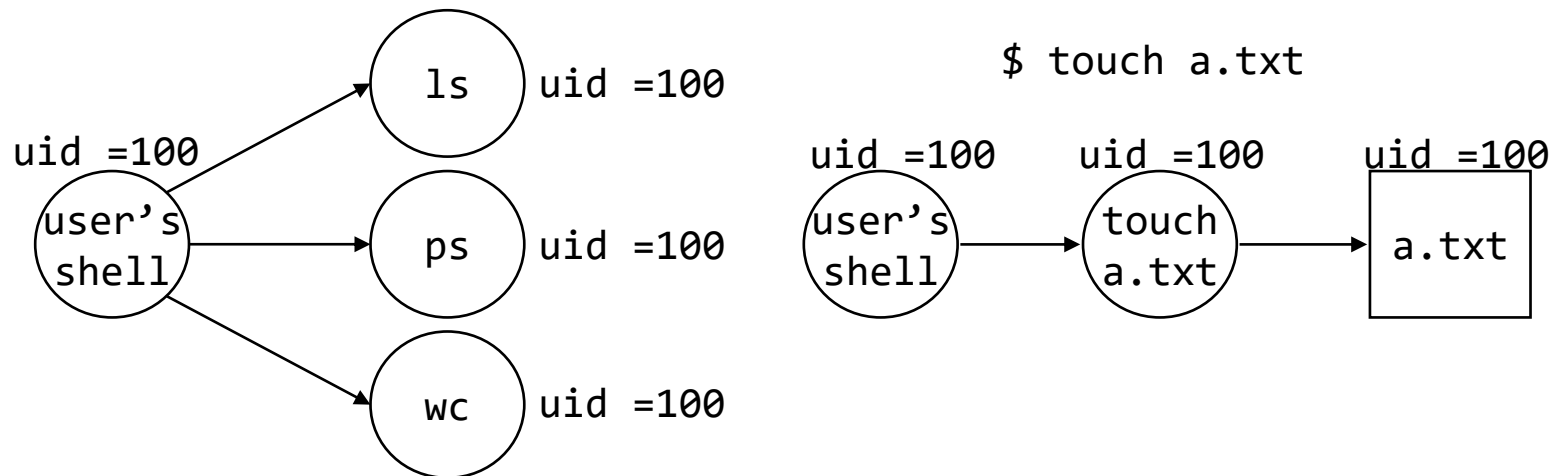
- Every file on Linux is owned by one of the system's users.
  - Normally the user who created the file.
- The owner's actual identity
  - `user-id (uid)`
  - `uid` associated with a particular username (in `/etc/passwd`)

```
sunwoo:x:1012:1012:,,,:/home/sunwoo:/bin/bash
```

<code>login-id:</code>	2-8 characters or digits
<code>password:</code>	encoded password of length 13
<code>uid:</code>	numeric user id between 0 and 60,000
<code>gid:</code>	numeric group id between 0 and 60,000
<code>User info:</code>	Actual user's name
<code>home-dir:</code>	user's home directory name
<code>shell:</code>	the default shell

# Users and Ownership (2/3)

- Each Unix / LINUX process is normally associated with `uid` of the user who launched the process.
- When a file is created by a process, the system establishes the ownership by referring to the `uid` of the process.
- Only the owner or superuser can change the file ownership
  - superuser (username = root, `uid` = 0)



# Users and Ownership (2/3)


- Group
  - Each user belongs to at least one group, possibly more.
  - Defined in `/etc/group`
  - Group identity
    - group-id (`gid`)
    - `gid` is associated with a particular group name.
- All files inherits `gid` and `uid` from the user process!

# Effective User and Group IDs

- Real user ID (`ruid`)
  - The current user's ID.
- Effective user ID (`euid`)
  - The `uid` of the user who evaluate privileges of the process to perform a particular action.
- In most cases, `ruid` and `euid` are the same.
- Otherwise, `euid` and `egid` determines file access permission.

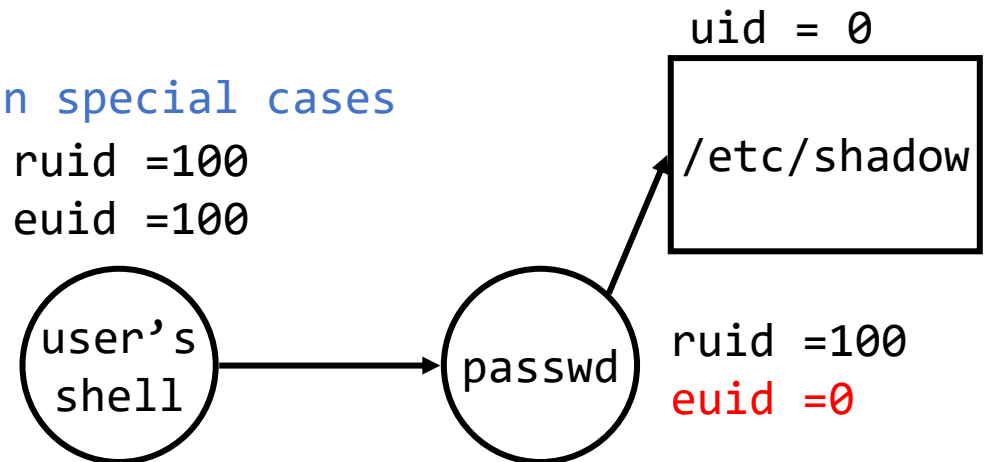
in most cases

`ruid =100`  
`euid =100`




in special cases

`ruid =100`  
`euid =100`



# Permission and File Modes (1/3)

- Ownership: the access permission associated with the target file.
- Permission
  - The right to read/write/execute the file.
  - Described as three-digit octal value.





# Permission and File Modes (2/3)

```
$ ls -l "filename"
```

```
$ touch test
```

```
$ ls -l test
```

```
-rw-r--r-- 1 s221394 s221394 0 Aug 15 14:58 test
```

↑ ↑ ↑ ↑  
File type user (owner) group others

# Permission and File Modes (3/3)

- Defined in <sys/stat.h>

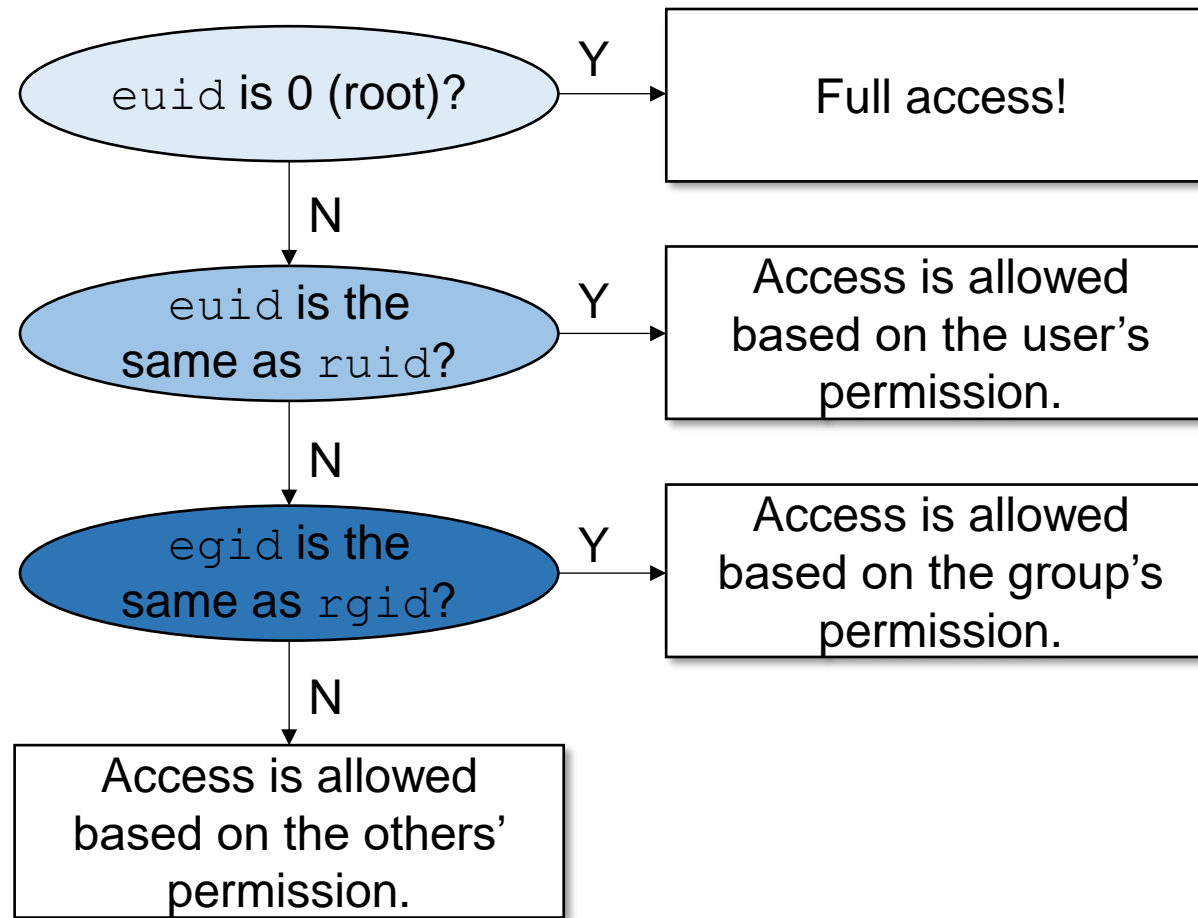
Octal value	Symbol	Permission
00400	S_IRUSR	read by owner
00200	S_IWUSR	write by owner
00100	S_IXUSR	execute by owner
00040	S_IRGRP	read by group
00020	S_IWGRP	write by group
00010	S_IXGRP	execute by group
00004	S_IROTH	read by others
00002	S_IWOTH	write by others
00001	S_IXOT	execute by others

S\_IRUSR | S\_IWUSR | S\_IXUSR | S\_IRGRP | S\_IXGRP | S\_IROTH | S\_IXOTH  
= 0755 = rwx-r-xr-x

# File Permission: `open()`

- When `open()` is called to access an existing file:
  - The kernel performs an access test using the effective user and group IDs (`eu`id and `eg`id).
  - If the process does not have the requested access permission, `open()` returns -1 (`errno` = `EACCESS`).

# File Access Test



# The File Creation Mask

- Each process is associated with a file creation mask.
  - The mask turns off particular permission bits whenever a file is created.
- The following two lines identically work.

```
filedes = open(pathname, O_WRONLY|O_CREAT, mode);  
filedes = open(pathname, O_WRONLY|O_CREAT, (~mask) & mode);
```

	← owner →			← group →			← others →		
	r	w	x	r	w	x	r	w	x
mode	1	1	1	1	1	1	1	1	1
mask	0	0	0	0	1	0	0	1	0
(~mask) & mode	1	1	1	1	0	1	1	0	1

# File Permission: `umask()`

- `umask()` sets the file creation mask.

```
# include <sys/stat.h>

mode_t umask(mode_t cmask);
```

```
mode_t oldmask;           /* to store the old mask value */

oldmask = umask(022);     /* block the write permission of others except for the owner */
```

***After the call of `umask()` above, any other processes do not have write permission on the current user's files.***

# File Permission: Example

```
#include <fcntl.h>
#include <sys/stat.h>

int specialcreat (const char *pathname, mode_t mode) {
    mode_t oldu;
    int filedes;

    /* set file creation mask as 0 */
    if ( (oldu = umask(0)) == -1) {
        perror ("saving old mask");
        return (-1);
    }

    /* create file */
    if((filedes=open(pathname, O_WRONLY | O_CREAT | O_EXCL, mode)) == -1)
        perror ("opening file");

    /* restore old mask even if the opening fails */
    if (umask (oldu) == -1)
        perror ("restoring old mask");

    return filedes;
}
```

# System Call: Access ()

```
# include <unistd.h>

int access(const char *pathname, int amode);
```

- `access ()` checks the permissions of `*pathname` using `ruid` and `rgid`.
- Arguments
  - `amode`
    - `R_OK`: test for read permission
    - `W_OK`: test for write permission
    - `X_OK`: test for execute permission
    - `F_OK`: test for existence of file



# Example of Access ( )

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

main()
{
    char *filename = "afile";

    if (access (filename, R_OK) == -1)
    {
        fprintf (stderr, "User cannot read file %s\n", filename);
        exit (1);
    }

    printf ("%s readable, proceeding\n", filename);

    ... /* the rest of the program */
}
```

# System Call: chmod ( )

```
# include <unistd.h>

int chmod(const char *pathname, mode_t newmode);
```

- chmod ( ) changes the permission of the file referenced to by pathname.

```
if (chmod(pathname, 0777) == -1)
    perror("call to chmod failed");
```

```
sunwool@lambda-server-vital:~/inha/tf2-fl$ chmod 0777 acc.txt
sunwool@lambda-server-vital:~/inha/tf2-fl$ ls -al acc.txt
-rwxrwxrwx 1 sunwool sunwool 479 Sep 15 23:58 acc.txt
sunwool@lambda-server-vital:~/inha/tf2-fl$
```

# System Call: chown ( )

```
# include <unistd.h>

int chown(const char *pathname, uid_t owner_id, gid_t group_id);
```

- chown ( ) changes the uid and gid of a file.
- Only the owner or root can change the owner of a file.
- EPERM error is returned when others change the owner.

- Files in Multi-user Environment
- ***Files with Multiple Names***
- Obtaining File Information

# Hierarchical File Sharing

## Process table

Process state
Process ID
User ID, group ID
Program file
File descriptor table
Memory mapping
Saved registers
Stack pointer

## File tables

File status flags
Current file offset
v-node pointer

File status flags
Current file offset
v-node pointer

## v-node tables

v-node information
i-node information
Current file size

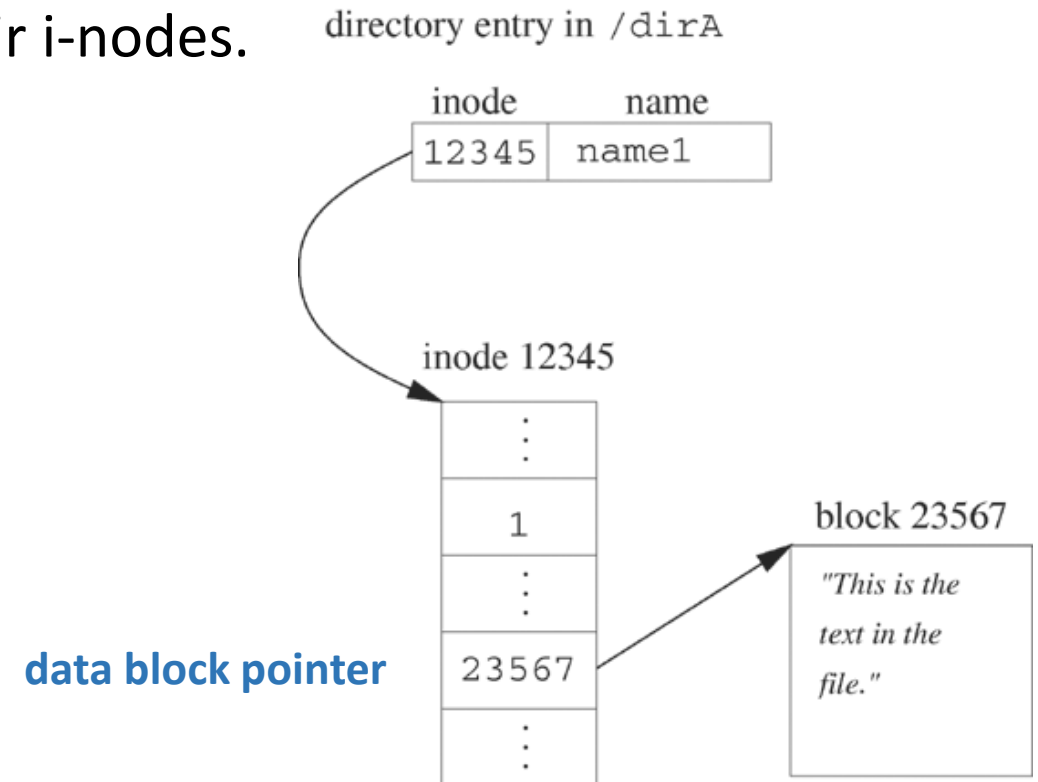
v-node information
i-node information
Current file size

file table  
entry pointer

fd flags	
fd 0	FD_CLOEXEC
fd 1	
fd 2	
...	

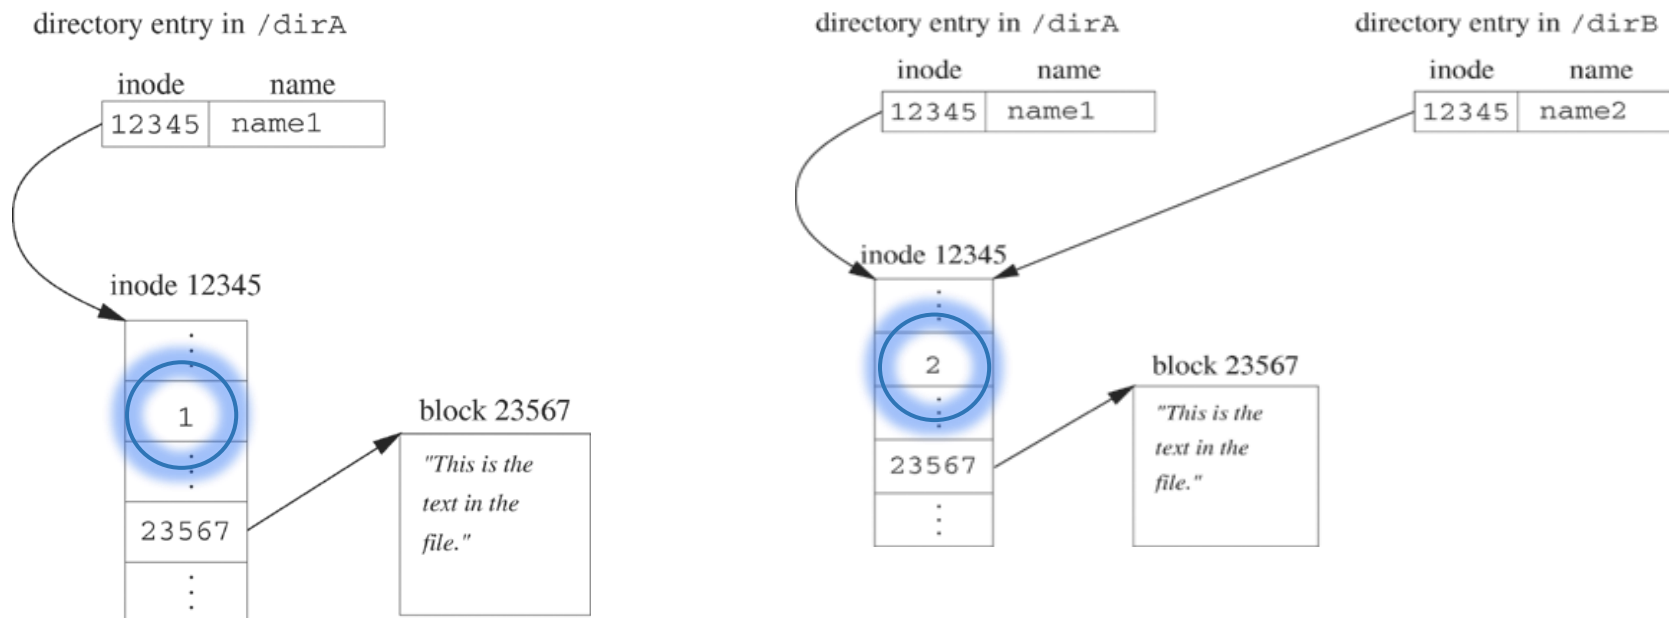
# i-node

- Each file has its own i-node.
  - Files have data blocks pointed from their i-nodes.
- i-node 0, 1, and 2 are pre-occupied.
  - 0: used to mean “no-i-node”
  - 1: used to collect bad disk blocks
  - 2: reserved for the root directory



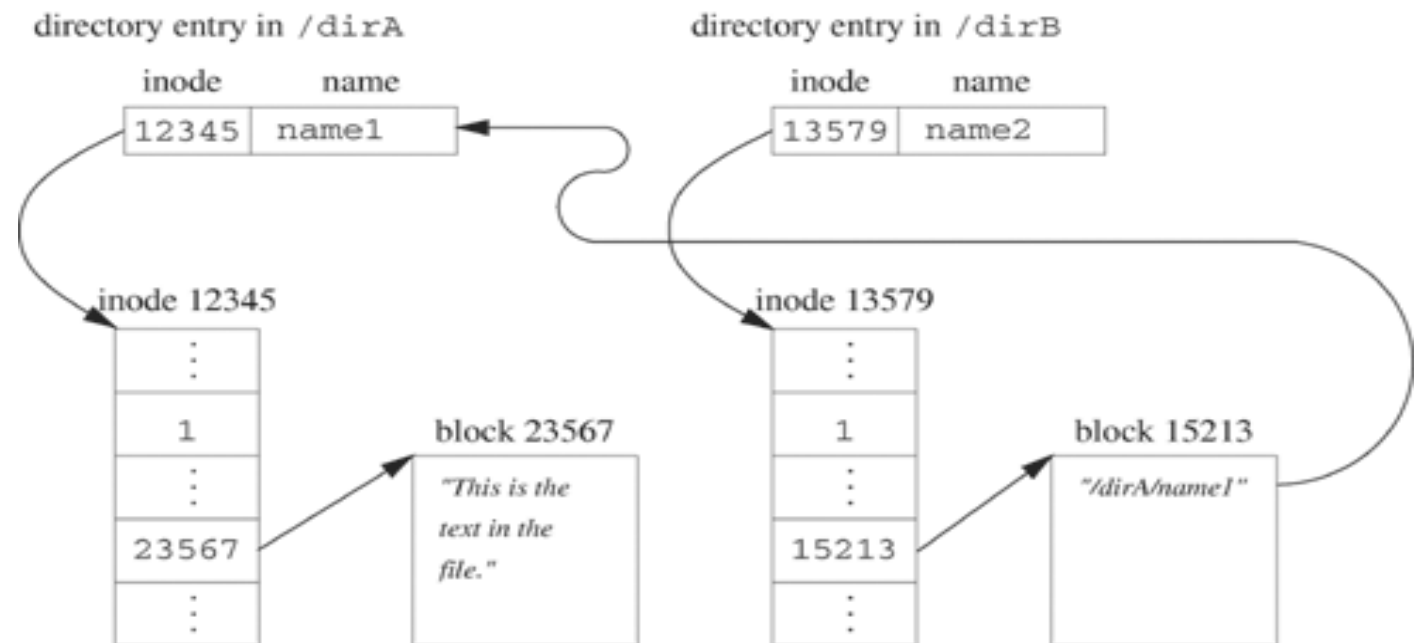
# Hard Link and Link Count

- Hard link: a direct pointer to a file
- Link count: the number of directory entries that point to the i-node
  - The actual block is deleted only when the link count is 0.
  - That is, even if the original file is deleted, the hard link file remains with the original block!



# Symbolic Link

- Symbolic link: an indirect pointer to a file
- Actual contents of a symbolic link is the file that is linked.
  - If the original file is deleted, the symbolic link becomes invalid.





# System Call: link()

```
# include <unistd.h>

int link(const char *original_path, const char *new_path);
```

- `link()` creates a new directory entry and increases the link count by one.
- Only superusers can create a hard link to directories.
- Both the original file and the new hard link file should be on the *same file system!*

# System Call: unlink()

```
# include <unistd.h>

int unlink(const char *pathname);
```

- `unlink()` removes an existing directory entry.
- It removes the link and reduces the file's link count by one.
  - If the link count is larger than 0, the data block is not returned.
  - If the link count becomes 0, the blocks are returned to the free block list.

# link() and unlink() Example

```
/* move - move a file from path name of file1 to file2 */
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>

char *usage = "usage: move file1 file2\n";

main (int argc, char **argv) {
    if (argc != 3){
        fprintf (stderr, usage); exit (1);
    }

    if ( link (argv[1], argv[2]) == -1) {
        perror ("link failed");
        exit (1);
    }

    if ( unlink(argv[1]) == -1) {
        perror ("unlink failed");
        unlink (argv[2]);
        exit (1);
    }
    printf ("Succeeded\n");
}
```

# Additional System Calls: remove()

```
# include <stdio.h>

int remove(const char *pathname);
```

- For files, `remove()` is identical to `unlink()`.
- `remove()` is ISO C standard while `unlink()` is Unix (LINUX) specific.

ISO C is a C language standard published by American National Standard Institute (ANSI).

# Additional System Calls: rename()

```
# include <stdio.h>

int rename(const char *oldname, const char *newname);
```

- `rename()` renames a file or a directory.
- In the ISO C standard, `rename()` works only for files not directories.

# Additional System Calls: symlink()

```
# include <unistd.h>

int symlink(const char *realname, const char *symname);
```

- `symlink()` makes a new file `symname` that points to the file `realname`.
- `open()` system call to the path of `symname` follows the path to `realname`.
- `readlink()` gives you the linked data of the file `symname`.

# Additional System Calls: readlink()

```
# include <unistd.h>

int readlink(const char *sympath, char *buffer, size_t bufsize);
```

- `readlink()` opens `sympath` file, reads the linked data into `buffer`, and closes the `sympath` file.
  - Like `read()` system call, `readlink()` returns the number of bytes successfully read.

- Files in Multi-user Environment
- Files with Multiple Names
- ***Obtaining File Information***




# System Call: stat() (1/2)

```
# include <sys/stat.h>

int stat(const char *pathname, struct stat *buf);
int fstat(int filedes, struct stat *buf);
int lstat(const char *pathname, struct stat *buf);
```

- These family functions read the file information and store it to the provided buffer.
  - `stat()`: file with a pathname.
  - `fstat()`: opened file with `filedes`
  - `lstat()`: symbolic link with `pathname`.



# System Call: stat() (2/2)

```
struct stat {  
    mode_t      st_mode;      /* file type & mode (permissions) */  
    ino_t       st_ino;       /* i-node number (serial number) */  
    dev_t       st_dev;       /* device number (file system) */  
    dev_t       st_rdev;      /* device number for special files */  
    nlink_t     st_nlink;     /* number of links */  
    uid_t       st_uid;       /* user ID of owner */  
    gid_t       st_gid;       /* group ID of owner */  
    off_t       st_size;      /* size in bytes, for regular files */  
    time_t      st_atime;     /* time of last access */  
    time_t      st_mtime;     /* time of last modification */  
    time_t      st_ctime;     /* time of last file status change */  
    blksize_t   st_blksize;   /* best I/O block size */  
    blkcnt_t    st_blocks;    /* number of disk blocks allocated */  
};
```

# Example: filedata

## (85 page of the textbook)

```
#include <stdio.h>
#include <sys/stat.h>

/* octarray is used to check if we have a proper
 * access permission. */
static short octarray[9] = {0400, 0200, 0100,
                           0040, 0020, 0010,
                           0004, 0002, 0001};

/* The length of perms should be 10 due to the NULL
 * at the end of the permission string. */
static char perms[10] = "rwxrwxrwx";

int filedata (const char *pathname)
{
    struct stat statbuf;
    char descrip[10];
    int j;

    if(stat (pathname, &statbuf) == -1) {
        fprintf (stderr, "Couldn't stat %s\n", pathname);
        return (-1);
    }

    /* Check all individual bit. */
    for(j=0; j<9; j++)
    {
        if (statbuf.st_mode & octarray[j])
            descrip[j] = perms[j];
        else
            descrip[j] = '-';
    }

    descrip[9] = '\\0';

    /* Print out the status of the file we collected. */

    printf ( " \\nFile %s :\\n", pathname);
    printf ( " Size %ld bytes\\n", statbuf.st_size);
    printf ( " User-id %d, Group-id %d\\n\\n",
            statbuf.st_uid, statbuf.st_gid);
    printf ( " Permissions: %s\\n", descrip);
    return (0);
}
```

# Example: lookout

*(76 page of the textbook)*

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/stat.h>

#define MFILE      10

void cmp(const char *, time_t);
struct stat sb;

main (int argc, char **argv)
{
    int j;
    time_t last_time[MFILE+1];

    if(argc < 2) {
        fprintf (stderr, "usage: lookout filename ...\n");
        exit (1);
    }

    if(--argc > MFILE) {
        fprintf (stderr, "lookout: too many filenames\n");
        exit (1);
    }
}
```

```
for (j=1; j<=argc; j++) { /* initialization */
    if (stat (argv[j], &sb) == -1) {
        fprintf (stderr,
            "lookout: couldn't stat %s\n", argv[j]);
        exit (1);
    }
    last_time[j] = sb.st_mtime;
}

for (;;) { /* loop until the file is changed. */
    for (j=1; j<=argc; j++)
        cmp (argv[j], last_time[j]);

    sleep (60);
}

void cmp(const char *name, time_t last){
    /* Check when the file was updated. */
    if (stat(name, &sb) == -1 || sb.st_mtime != last) {
        fprintf (stderr, "lookout: %s changed\n", name);
        exit (0);
    }
}
```

# Wrap-Up

- **Ownership:** Whenever a process accesses a file, it performs the permission test to check the file ownership.
  - Effective user ID (`euuid`) should be the same as real user ID (`ruuid`) to modify or delete the file.
- **File link:**
  - Hard link file points to the same i-node as the original file.
  - Symbolic link has its own i-node which points again the original file.
- **File information:**
  - `stat` family functions extract the information using the file name, the file descriptor, or the symbolic link.

# Any Questions?

